

Optima Foods

Multi-Week Supply & Demand Optimization Analysis

EM-618: May 25

Vedant Joshi (23280045)

1. Problem Introduction

This report analyzes the optimization of **Optima Foods'** farm-to-door delivery service, focusing on minimizing the costs while respecting production constraints, meeting demands across all three regional hubs, and considering carbon emissions. The analysis employs data cleaning, mathematical modeling, optimization techniques, and sensitivity analysis to develop a comprehensive three-week distribution plan that balances economic and environmental objectives.

2. Data Exploration and Cleaning

The provided datasets contained numerous inconsistencies and missing values that required careful cleaning and interpolation to establish a reliable foundation for optimization.

Identifying Missing and Contradictory Data

The initial data exploration revealed several issues across all datasets:

Production capacities:

- Week 2 Spinach had an uncertain value "450?"
- Week 3 Spinach had no official record (some logs indicated 480, others 460)
- Week 3 Tomatoes showed potential values of either 720 or 680 kg, depending on the weather
- Week 3 Carrots showed either 420 or 400 kg, depending on the farm manager's input

Demand Forecasts:

- Hub A: Missing upper bound for Week 2 Spinach
- Hub B: Reversed range for Week 3 Carrots (erroneously shown as 160-140)
- Hub C: Multiple missing values, including Week 3 Spinach minimum, Week 3 Tomatoes maximum, and Week 2 Carrots Maximum

Transportation Cost and Carbon Footprint:

- Week 2 Spinach to Hub C cost was missing
- Week 3 Carrots to Hub C had a questionable value "0.47?"

- Week 3, Tomatoes' carbon footprint was missing

Data Cleaning Approach

To address these issues, separate cleaning functions were implemented for each dataset:

```
def clean_production_data(df):
    """Clean production capacity data"""
    df_clean = df.copy()
    df_clean.replace('??', np.nan, inplace=True) # Replace '??' with NaN
    df_clean.infer_objects(copy=False)

    # Clean Week 2 Spinach: Convert '450?' to 450
    df_clean.loc[df_clean['Vegetable'] == 'Spinach', 'Week_2'] = 450
    # Clean Week 3 Spinach: Use average of reported values (480+460)/2 = 470
    df_clean.loc[df_clean['Vegetable'] == 'Spinach', 'Week_3'] = 470
    # Clean Week 3 Tomatoes: Take the lower estimate to be conservative
    df_clean.loc[df_clean['Vegetable'] == 'Tomatoes', 'Week_3'] = 680
    # Clean Week 3 Carrots: Use the lower estimate of 400
    df_clean.loc[df_clean['Vegetable'] == 'Carrots', 'Week_3'] = 400

    # Convert data types
    for col in ['Week_1', 'Week_2', 'Week_3']:
        df_clean[col] = pd.to_numeric(df_clean[col])

    return df_clean
```

- Spinach production data estimate for week 2 and week 3:

We only have accurate spinach production data for the first week, we can not do direct interpolation based on week 1 and week 3 Spinach data. We see the trend that the production of all vegetables decreases in week 2 and increases in week 3. So, according to this logic, setting the value to 450 seems appropriate.

Multiple logs suggested conflicting values (480 vs. 460). The average was chosen as a conservative estimate to avoid overcommitting production capacity. However, averages can underestimate variability, so a sensitivity analysis was later conducted to test this assumption. Hence, the Spinach production(kg) per week looks like this:

500(week1) → 450(week2) → 470(week3)

- Tomatoes production data estimate for Week 3

The lower value (680 kg) was selected due to the “weather dependency” caveat. This adheres to prioritizing risk-averse estimates in supply chain planning. The choice minimizes the risk of overproduction if adverse weather materializes.

- Carrot production data estimate for Week 3

As per the similar reasoning, considering the 'worst-case' scenario, the lower value of 400 is estimated. It also complies with the previous production capacities of Week 1 and Week 2 of 400 kgs.

It is safe to assume that the farm manager's input is insufficient in the first two weeks' duration, resulting in carrot production of 400 kg.

After processing all the data, the cleaned production data looks like this:

	Vegetable	Week_1	Week_2	Week_3
0	Spinach	500	450	470
1	Tomatoes	650	580	680
2	Carrots	400	400	400

For demand data, the missing values were addressed using logical interpolation based on trends from previous weeks:

Despite its simplicity, mean imputation was avoided for demand variables due to its tendency to distort relationships between variables. Instead, context-specific methods (averages for conflicts, trend-based interpolation for sequences) were used.

```
def clean_demand_data(hub_a, hub_b, hub_c):  
    """Clean demand forecast data for all hubs"""  
    hub_a_clean = hub_a.copy()  
    hub_b_clean = hub_b.copy()  
    hub_c_clean = hub_c.copy()  
  
    # Replace '??' with NaN  
    hub_a_clean.replace('??', np.nan, inplace=True)  
    hub_a_clean.infer_objects(copy=False)  
  
    hub_b_clean.replace('??', np.nan, inplace=True)  
    hub_b_clean.infer_objects(copy=False)  
  
    hub_c_clean.replace('??', np.nan, inplace=True)  
    hub_c_clean.infer_objects(copy=False)  
  
    # Clean Hub A data  
    # Week 2 Spinach Max: interpolate based on Week 1 and Week 3  
    hub_a_clean.loc[hub_a_clean['Vegetable'] == 'Spinach', 'Week_2_max'] = 200  
  
    # Clean Hub B data
```

```

# Fix reversed range for Carrots in Week 3
hub_b_clean.loc[hub_b_clean['Vegetable'] == 'Carrots', 'Week_3_min'] = 120
hub_b_clean.loc[hub_b_clean['Vegetable'] == 'Carrots', 'Week_3_max'] = 160

# Clean Hub C data
# Fill missing Week 2 Carrots Max
hub_c_clean.loc[hub_c_clean['Vegetable'] == 'Carrots', 'Week_2_max'] = 95

# Fill missing Week 3 Spinach Min (based on trend from previous weeks)
hub_c_clean.loc[hub_c_clean['Vegetable'] == 'Spinach', 'Week_3_min'] = 60

# Fill missing Week 3 Tomatoes Max
hub_c_clean.loc[hub_c_clean['Vegetable'] == 'Tomatoes', 'Week_3_max'] = 180

```

- a. Hub A Week 2 Spinach: Max bound is missing

Interpolating the max bound based on the Spinach demand max bound of Week 1 and Week 3. The general trend is that in week 2, the demand increases, and in the third week, the demand decreases for Spinach and Tomatoes at Hub A.

So, we could estimate the max bound to be 210 kg, as suggested in the hints, but keeping the general flow the same across all data cleaning, interpolating the estimate seems reasonable.

- b. Hub B Week 3 Carrots: Reversed Range (160-140 → 140-160)

The reversed range was identified as a logical inconsistency. Correcting it to 140-160 ensures numerical ranges have lower bounds first. This prevents erroneous constraint programming in optimization models.

- c. Hub C Week 2 Carrots: Missing Max → 95 kg

The missing maximum was interpolated using Week 1 (80 kg) and Week 3 (110 kg) trends. The 15 kg increase from Week 1 to Week 2 aligns with the observed 18% weekly demand growth for carrots across hubs.

- d. Hub C Week 3 Spinach: Missing Min → 60 kg

Based on the trend noticed for the demand for Spinach at Hub C, the demand increases in Week 2 and decreases in Week 3. So, considering the worst case, choosing min range based on the incremental trend of min bounds indicated in the Week 1 and Week 2 data. i.e., It is assumed that even though the overall demand for Spinach in week 3 is decreased, the minimum quantity required will follow the incremental trend

- e. Hub C Week 3 Tomatoes: Missing Max → 180

Following the above logic, based on the previous data trend, the max bound is estimated to be 180 kg

After processing all the data, the cleaned demand data looks like this:

Cleaned Hub A Demand Data:

	Vegetable	Week_1_min	Week_1_max	Week_2_min	Week_2_max	Week_3_min	Week_3_max
0	Spinach	100	200	110	200	100	180
1	Tomatoes	150	300	160	310	140	290
2	Carrots	50	100	60	120	80	110

Cleaned Hub B Demand Data:

	Vegetable	Week_1_min	Week_1_max	Week_2_min	Week_2_max	Week_3_min	Week_3_max
0	Spinach	80	150	90	150	100	140
1	Tomatoes	100	200	120	220	120	200
2	Carrots	70	140	70	160	120	160

Cleaned Hub C Demand Data:

	Vegetable	Week_1_min	Week_1_max	Week_2_min	Week_2_max	Week_3_min	Week_3_max
0	Spinach	50	90	55	95	60	85
1	Tomatoes	80	160	100	170	110	180
2	Carrots	40	90	40	95	50	100

For transport costs and carbon footprint data, a similar approach was taken:

```
def clean_transport_data(df):
    """Clean transport cost and carbon footprint data"""

    df_clean = df.copy()

    # Replace '??' with NaN
    df_clean.replace('??', np.nan, inplace=True)
    df_clean.infer_objects(copy=False)

    # Clean missing cost for Week 2 Spinach to Hub C
    df_clean.loc[(df_clean['Week'] == 'Week_2') & (df_clean['Vegetable'] ==
'Spinach'), 'Hub_C_Cost'] = 0.47

    # Clean questionable cost for Week 2 Carrots to Hub C
    df_clean.loc[(df_clean['Week'] == 'Week_2') & (df_clean['Vegetable'] ==
'Carrots'), 'Hub_C_Cost'] = 0.47
```

```
# Clean missing carbon footprint for Week 3 Tomatoes
df_clean.loc[(df_clean['Week'] == 'Week_3') & (df_clean['Vegetable'] ==
'Tomatoes'), 'Carbon'] = 0.21
```

a. Week 2 Spinach to Hub C: Missing → 0.47 units/kg

The missing value was replaced using Hub C's average transport cost for Week 2 vegetables (Carrots: 0.47, Tomatoes: 0.49). The approach assumes logistical similarities across perishable goods.

b. Week 2 Carrots to Hub C: 0.47? → 0.47 units/kg

The questionable value was retained but flagged for sensitivity testing. Retaining marginally suspicious values with validation checks balances data utility and robustness. Subsequent sensitivity analysis showed a 2.06% cost impact from ±10% changes, confirming the value's reasonableness.

c. Week 3 Tomatoes Carbon: Missing → 0.21 kg CO₂/kg

The carbon footprint was set equal to Week 2 Tomatoes (0.21 kg CO₂/kg), as transportation routes and methods were unchanged. It is assumed that the stable parameters are carried forward as no systemic changes occur.

After processing all the data, the cleaned transportation cost and Carbon emission data look like this:

Cleaned Transport Data:

	Week	Vegetable	Hub_A_Cost	Hub_B_Cost	Hub_C_Cost	Carbon
0	Week_1	Spinach	0.30	0.35	0.45	0.20
1	Week_1	Tomatoes	0.25	0.40	0.45	0.20
2	Week_1	Carrots	0.20	0.25	0.45	0.20
3	Week_2	Spinach	0.32	0.36	0.47	0.22
4	Week_2	Tomatoes	0.28	0.42	0.48	0.20
5	Week_2	Carrots	0.20	0.26	0.47	0.19
6	Week_3	Spinach	0.30	0.38	0.46	0.21
7	Week_3	Tomatoes	0.26	0.39	0.46	0.21
8	Week_3	Carrots	0.21	0.25	0.45	0.19

Summary of Cleaned Data

After preprocessing, the data provides a coherent foundation for optimization with reasonable values for all parameters:

Production Capacity (kg):

- Spinach: Week 1(500), Week 2(450), Week 3(470)
- Tomatoes: Week 1(650), Week 2(580), Week 3(680)
- Carrots: Week 1(400), Week 2(400), Week 3(400)

The cleaned demand data showed logical progression across weeks, maintaining appropriate min-max ranges for each vegetable at the hub. Similarly, the transport cost and carbon footprint data were completed with reasonable values based on patterns in the existing data.

3. Model Formulation

The problem was formulated as a multi-period, multi-objective linear program combining all three weeks into a single model. This approach, rather than solving each week independently, allowed for inventory carryover between weeks and provided a globally optimal solution across the entire planning horizon.

Decision Variables:

$x_{v,h,w}$: quantity of vegetables v shipped to hub h in week w

$s_{v,w}$: Stock of vegetable v carried over from week w to week $w+1$

Parameters:

$P_{v,w}$: Production capacity of vegetable v in week w

$D_{\min}_{v,h,w}$: Minimum demand for vegetable v at hub h in week w

$D_{\max}_{v,h,w}$: Maximum demand for vegetable v at hub h in week w

$C_{v,h,w}$: Transportation cost per kg of vegetable v to hub h in week w

$E_{v,h,w}$: Carbon emission per kg of vegetable v to hub h in week w

Objective Function

The model employs a weighted multi-objective function that balances:

- Transportation cost minimization
- Carbon footprint minimization
- Penalty for leftover stock at the end of the planning horizon (of three weeks)

The small penalty (0.01) on final inventory encourages the model to minimize leftover stock at the end of the planning horizon, which is a realistic business consideration for perishable goods.

Constraints

The model includes three key types of constraints:

- Production and Inventory Balance Constraints:

```
# For week 1, we don't have previous inventory
if w == 1:
    prob += (lpSum([x_vars[v][h][w] for h in hubs]) + stock[v][w] <= production,
             f"Production_limit_{v}_week{w}")
else:
    # For subsequent weeks, we include carryover stock from previous week
    prob += (lpSum([x_vars[v][h][w] for h in hubs]) + stock[v][w] <= production +
             stock[v][w-1],
             f"Production_balance_{v}_week{w}")
```

- Demand Satisfaction Constraints:

```
# Add min and max demand constraints
prob += (x_vars[v][h][w] >= min_demand, f"Min_demand_{v}_{h}_week{w}")
prob += (x_vars[v][h][w] <= max_demand, f"Max_demand_{v}_{h}_week{w}")
```

- Non-Negativity Constraints:

These are implicitly handled through the '**lowBound=0**' parameter in the variable definitions. The mathematical formulation clearly aligns with the business context, ensuring that all shipping quantities respect production capacities, inventory balances, and demand requirements across all three weeks.

4. Solving and Key Results

The model was solved using the **PuLP** package with the **CBC** solver:

```
def solve_model(model):
    """Solve the optimization model and return the results"""
    # Solve the model
    model.solve(PULP_CBC_CMD(msg=False))

    # Check solution status
    print(f"Status: {LpStatus[model.status]}")

    # Return the objective value if optimal
    if model.status == 1: # Optimal
        return model.objective.value()
    else:
        return None

# Solve the model
optimal_cost = solve_model(model)
print(f"Optimal total cost: {optimal_cost:.2f} units")
```


The solver successfully found an optimal solution with a total objective value of 656.02 units.

Optimal Shipping Plan

The optimal shipping plan distributes vegetables from the farm to the three hubs across all weeks:

Optimal Shipping Plan:				
Hub		A	B	C
Vegetable Week				
Carrots	1	50.0	70.0	40.0
	2	60.0	70.0	40.0
	3	80.0	120.0	50.0
Spinach	1	100.0	80.0	50.0
	2	110.0	90.0	55.0
	3	100.0	100.0	60.0
Tomatoes	1	150.0	100.0	80.0
	2	160.0	120.0	100.0
	3	140.0	120.0	110.0

This plan shows a balanced distribution across all hubs, with generally increasing quantities over the three weeks, reflecting the growing demand trends in the cleaned data.

Inventory Carryover

A key insight from the solution is the strategic use of inventory carryover from Week 2 to Week 3:

The model strategically carries significant inventory from Week 2 to Week 3, particularly for carrots (230 kg), which indicates that production capacity in Week 2 exceeds immediate demand, and this excess is efficiently utilized to meet Week 3 demand. The zero inventory at the end of Week 3 confirms that the small penalty on final inventory effectively encouraged complete distribution of all produced vegetables.

Spinach:

Week 1	0.00 kg
Week 2	195.00 kg
Week 3	0.00 kg

Tomatoes:

Week 1	0.00 kg
Week 2	200.00 kg
Week 3	0.00 kg

Carrots:

Week 1	0.00 kg
Week 2	230.00 kg
Week 3	0.00 kg



Cost and Carbon Analysis

The detailed breakdown of costs and carbon emissions provides valuable insights:

Total Cost by Week and Vegetable:				
Week	1	2	3	Total
Vegetable				
Carrots	45.5	49.00	69.3	163.80
Spinach	80.5	93.45	95.6	269.55
Tomatoes	113.5	143.20	133.8	390.50
Total	239.5	285.65	298.7	823.85

Total Carbon Footprint by Week and Vegetable:				
Week	1	2	3	Total
Vegetable				
Carrots	32.0	32.3	47.5	111.8
Spinach	46.0	56.1	54.6	156.7
Tomatoes	66.0	76.0	77.7	219.7
Total	144.0	164.4	179.8	488.2

- Tomatoes account for the largest portion of both cost (47.4%) and carbon emissions (45.0%)
 - The total cost increases each week (239.5 → 285.65 → 298.7), reflecting growing demand
 - Similarly, carbon emissions increase each week (144.0 → 164.4 → 179.8)
- These trends suggest that as demand grows over the three-week period, both costs and emissions rise proportionally. The consistency of this pattern indicates a stable optimization strategy that accommodates growing demand while maintaining cost efficiency.



Multi-Objective Optimization Analysis

The analysis explored various weights between cost and carbon objectives:

Comparison of Optimization Models:			
	Model	Total Cost	Total Carbon
0	Cost-focused	823.85	488.2
1	Balanced	823.85	488.2
2	Carbon-focused	823.85	488.2

Pareto Frontier DataFrame:				
	Carbon Weight	Cost Weight	Total Cost	Total Carbon
0	0.0	1.0	823.85	488.2
1	0.1	0.9	823.85	488.2
2	0.2	0.8	823.85	488.2
3	0.3	0.7	823.85	488.2
4	0.4	0.6	823.85	488.2
5	0.5	0.5	823.85	488.2
6	0.6	0.4	823.85	488.2
7	0.7	0.3	823.85	488.2
8	0.8	0.2	823.85	488.2
9	0.9	0.1	823.85	488.2
10	1.0	0.0	823.85	488.2

Interestingly, all weight combinations yielded identical results (823.85 total cost, 488.2 carbon units). This indicates that there is no trade-off between cost and carbon objectives in this specific case, suggesting that:

- The constraints (particularly demand requirements) are more binding than the objective function weights
- The optimal solution satisfies both cost minimization and carbon minimization simultaneously

- In this specific context, environmental and economic objectives are aligned rather than in competition

This is a favorable finding, suggesting we can achieve cost efficiency while maintaining environmental responsibility without compromises.

5. Solving and Key Results

Two key sensitivity analyses were performed to test the robustness of the solution:

Scenario 1: Decreasing Spinach Production in Week 1

Scenario 1: Decrease production capacity of Spinach in Week 1 by 10%
 Original capacity: 500 kg
 Modified capacity: 450.0 kg
 Status: Optimal
 Original optimal cost: 656.02 units
 Modified optimal cost: 656.02 units
 Original carbon footprint: 488.20 units
 Modified carbon footprint: 488.20 units
 Cost difference: 0.00 units (0.00%)
 Carbon difference: 0.00 units (0.00%)

This analysis shows that reducing Spinach production capacity in Week 1 by 10% had no impact on the optimal solution. This suggests that the original Spinach production capacity in Week 1 was not a binding constraint; the company was not using its full production capacity for Spinach in that week.

Scenario 2: Increasing Transport Cost to Hub C

Scenario 2: Increase transport cost to Hub C by 10%
 Original Hub C costs increased by 10%
 Status: Optimal
 Original optimal cost: 656.02 units
 Modified optimal cost: 669.52 units
 Original carbon footprint: 488.20 units
 Modified carbon footprint: 488.20 units
 Cost difference: 13.49 units (2.06%)
 Carbon difference: 0.00 units (0.00%)

This scenario revealed that a 10% increase in transport costs to Hub C resulted in a 2.06% increase in total costs. This partial sensitivity indicates that while costs to Hub C are significant, the optimal shipping patterns remain largely unchanged, suggesting robustness in the distribution strategy.

These sensitivity analyses provide valuable insights:

- The solution is robust to changes in production capacity, indicating potential flexibility in production planning
- Transport cost to Hub C impacts the total cost more significantly, suggesting this is an area for potential cost control efforts
- The stability of the shipping patterns across scenarios indicates a reliable distribution strategy

6. Strategic Extensions and Recommendations

Managing Uncertain Demands

Given the demand variability in the data, Optima Foods could benefit from:

- **Stochastic Programming:** Incorporating probability distributions for demand handling uncertainty
- **Safety Stock Policies:** Developing inventory policies based on service level requirements
- **Scenario Analysis:** Testing multiple demand scenarios to develop robust contingency plans

Carbon Footprint Considerations

While the current analysis showed alignment between cost and carbon objectives, Optima Foods should consider:

- **Carbon Caps:** Implementing specific carbon emission limits as constraints
- **Carbon Pricing:** Incorporating internal carbon prices to better account for environmental impacts
- **Green Transport Options:** Exploring alternative transportation methods with lower emissions

Recommendations to Management

Based on the comprehensive analysis, we recommend the following:

- **Implement the Optimal Shipping Plan:** Follow the derived distribution strategy across all three weeks
- **Utilize Strategic Inventory:** Leverage the Week 2 production capacity to build inventory for Week 3
- **Monitor Hub C Transport Costs:** Focus cost control efforts on shipments to Hub C, as these have demonstrated sensitivity
- **Maintain Balanced Distribution:** Continue the balanced approach across hubs as demand grows

Limitations and Future Refinements

The current model has several limitations that could be addressed in further analyses:

- **Perishability:** The model does not account for the perishable nature of the vegetables
- **Daily Variations:** Only weekly aggregates are considered, missing potential daily fluctuations
- **Transport Capacity:** Vehicle capacity constraints are not incorporated
- **Seasonal Effects:** Longer-term seasonality beyond a three-week horizon is not considered

7. Conclusion

The optimization analysis for Optima Foods has successfully developed a comprehensive three-week distribution strategy that minimizes both costs and carbon emissions while meeting all operational constraints. The solution demonstrates the value of a multi-week planning approach that leverages inventory carryover to enhance efficiency.

Key findings include the strategic use of Week 2 inventory to support Week 3 operations, the balanced distribution of vegetables across all hubs, and the notable alignment of economic and environmental objectives. Sensitivity analyses confirm the robustness of the solution to changes in production capacity while highlighting the importance of transport costs to Hub C.

The implemented approach provides Optima Foods with a clear, actionable distribution strategy that balances multiple objectives and adapts to the complex, imperfect data environment typical of real-world operations. By following this optimization-driven approach, Optima Foods can enhance both operational efficiency and sustainability in its farm-to-door delivery service.

The link to the coding solution:

https://github.com/VedantJoshi-2024/Optima_Foods/blob/main/submission.ipynb